

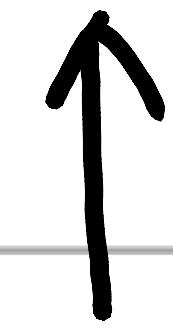
Go benchmark measures Cpu & memory metrics as a part of `StartTimer` & `StopTimer` methods.

These methods are invoked from `runN`, which runs N iterations of a benchmark.

```
func (b *B) runN(n int) {
    benchmarkLock.Lock()
    defer benchmarkLock.Unlock()
    ctx, cancelCtx := context.WithCancel(context.Background())
    defer func() {
        b.runCleanup(normalPanic)
        b.checkRaces()
    }()
    // Try to get a comparable environment for each run
    // by clearing garbage from previous runs.
    runtime.GC()
    b.resetRaces()
    b.N = n
    b.loop.n = 0
    b.loop.i = 0
    b.loop.done = false
    b.ctx = ctx
    b.cancelCtx = cancelCtx
    |
    b.parallelism = 1
    b.ResetTimer()
    b.StartTimer()
    b.benchFunc(b)
    b.StopTimer()
    b.previousN = n
    b.previousDuration = b.duration
```

Please note `RUNN` is called from  
`launch()` method.

```
func (b *B) launch() { → launch
    ...
    defer func() {
        b.signal = true
    }()
    ...
    if b.loop.n == 0 {
        // Run the benchmark for at least the specified amount of time.
        if b.benchTime.n > 0 {
            ...
            if b.benchTime.n > 1 {
                b.runN(b.benchTime.n)
            }
        } else {
            d := b.benchTime.d
            for n := int64(1); !b.failed && b.duration < d && n < 1e9; {
                last = n
                // Predict required iterations.
                goalns = d.Nanoseconds()
                prevIters := int64(b.N)
                n = int64(predictN(goalns, prevIters, b.duration.Nanoseconds(), last))
                b.runN(int(n))
            }
        }
    }
}
b.result = BenchmarkResult{ N: b.N, T: b.duration, Bytes: b.bytes, MemAllocs: b.netAllocs, MemBytes: b.netBytes, Extra: b.extra}
```



Store result

Let's cover the following:

- ① start timer
- ② stop timer
- ③ reset timer
- ④ BenchmarkResult
- ⑤ add result · (from Run(..))

```
func (b *B) StartTimer() {
    if !b.timerOn {
        runtime.ReadMemStats(&memStats)
        b.startAllocs = memStats.Mallocs
        b.startBytes = memStats.TotalAlloc
        b.start = highPrecisionTimeNow()
        b.timerOn = true
        b.loop.i &^= loopPoisonTimer
    }
}
```

StartTimer

time.Now()

(more on  
this here:  
)

- collect memory statistics

- capture start time



time.Now()

Memory statistics  
↳ TotalAlloc:  
cumulative bytes allocated

↳ Mallocs:  
↳ cumulative count  
of heap allocations

```
func (b *B) StopTimer() {
    if b.timerOn {
        b.duration += highPrecisionTimeSince(b.start)
        runtime.ReadMemStats(&memStats)
        b.netAllocs += memStats.Mallocs - b.startAllocs
        b.netBytes += memStats.TotalAlloc - b.startBytes
        b.timerOn = false
        // If we hit B.Loop with the timer stopped, fail.
        b.loop.i |= loopPoisonTimer
    }
}
```

## StopTimer

$\text{time.Now()} - \text{b.start}$

- Capture duration
- Capture allocation
- Compute net allocations & netBytes

```
func (b *B) ResetTimer() {
    if b.extra == nil {
        // Allocate the extra map before reading memory stats.
        // Pre-size it to make more allocation unlikely.
        b.extra = make(map[string]float64, 16)
    } else {
        clear(b.extra)
    }

    if b.timerOn {
        runtime.ReadMemStats(&memStats)
        b.startAllocs = memStats.Mallocs
        b.startBytes = memStats.TotalAlloc
        b.start = highPrecisionTimeNow()
    }

    b.duration = 0
    b.netAllocs = 0
    b.netBytes = 0
}
```

## ResetTimer

- Capture allocation, & time, if timer is on.
- Clear duration, net allocation & netBytes.

Everytime runN finishes,

properties like duration,

net allocation & netBytes are

set in the pointer to B.

(represents one  
benchmark)

After N iterations are done

for a benchmark, an instance

of BenchmarkResult is created.

```
func (b *B) launch() {
    b.signal <- true
}()
/...
if b.loop.n == 0 {
    // Run the benchmark for at least the specified amount of time.
    if b.benchTime.n > 0 {
        // We already ran a single iteration in run1.
        // If -benchtime=1x was requested, use that result.
        // See https://golang.org/issue/32051.
        if b.benchTime.n > 1 {
            b.runN(b.benchTime.n)
        }
    } else {
        d := b.benchTime.d
        for n := int64(1); !b.failed && b.duration < d && n < 1e9; {
            last := n
            // Predict required iterations.
            goalns := d.Nanoseconds()
            prevIters := int64(b.N)
            n = int64(predictN(goalns, prevIters, b.duration.Nanoseconds(), last))
            b.runN(int(n))
        }
    }
}
b.result = BenchmarkResult{ N: b.N, T: b.duration, Bytes: b.bytes, MemAllocs: b.netAllocs, MemBytes: b.netBytes, Extra: b.extra}
```



capture BenchmarkResult

after all iterations are done  
for a benchmark.

Effectively, b.N, b.duration,  
b.bytes, b.netAllocs, b.netBytes

are reported for the last iteration  
value.

```
type BenchmarkResult struct {
    N           int          // The number of iterations.
    T           time.Duration // The total time taken.
    Bytes       int64        // Bytes processed in one iteration.
    MemAllocs   uint64       // The total number of memory allocations.
    MemBytes    uint64       // The total number of bytes allocated.

    // Extra records additional metrics reported by ReportMetric.
    Extra map[string]float64
}
```

BenchmarkResult has N, T, MemAllocs and MemBytes.

```
func (b *B) Run(name string, f func(b *B)) bool {
    sub := &B{
        common: common{
            signal: make(chan bool),
            name:   benchName,
            parent: &b.common,
            level:  b.level + 1,
            creator: pc[:n],
            w:      b.w,
            chatty: b.chatty,
            bench:  true,
        },
        importPath: b.importPath,
        benchFunc:  f,
        benchTime:  b.benchTime,
        bstate:    b.bstate,
    }
    if partial {...}

    if b.chatty != nil {...}

    if sub.run1() {
        sub.run()
    }
    b.add(sub.result)
    return !sub.failed
}
```

After Run()

method is  
done for a

benchmark,

the add()

method is

called to

add the

Benchmark -

Result.

```
func (b *B) add(other BenchmarkResult) {
    r := &b.result
    // The aggregated BenchmarkResults resemble running all subbenchmarks as
    // in sequence in a single benchmark.
    r.N = 1
    r.T += time.Duration(other.NsPerOp())
    if other.Bytes == 0 {
        // Summing Bytes is meaningless in aggregate if not all subbenchmarks
        // set it.
        b.missingBytes = true
        r.Bytes = 0
    }
    if !b.missingBytes {
        r.Bytes += other.Bytes
    }
    r.MemAllocs += uint64(other.AllocsPerOp())
    r.MemBytes += uint64(other.AllocatedBytesPerOp())
}
```

The `add(..)` method simply computes  
`NsPerOp()`, `AllocsPerOp()`,  
`AllocatedBytesPerOp()` ...

\* `PerOp` represents per iteration.

```
func (r BenchmarkResult) AllocsPerOp() int64 {
    if v, ok := r.Extra["allocs/op"]; ok { return int64(v) }
    if r.N <= 0 { return 0 }
    return int64(r.MemAllocs) / int64(r.N)
}
```

The above code `AllocsPerOp()` is

simply memAllocs

N

> iterations

Other `PerOp()` methods are

similar.

Open:

Is the reported time

Also the diff.  
b/w the  
two

Cputime or usertime?