

Benchmark execution starts with finding all benchmarks to run.

```
size_t RunSpecifiedBenchmarks(BenchmarkReporter* display_reporter,
  if (!fname.empty()) {
    file_reporter->SetErrorStream(&output_file);
  }

  std::vector<internal::BenchmarkInstance> benchmarks;
  if (!FindBenchmarksInternal(spec, &benchmarks, &Err)) {
    Out.flush();
    Err.flush();
    return 0;
  }

  if (benchmarks.empty()) {
    Err << "Failed to match any benchmarks against regex: " << spec << "\n";
    Out.flush();
    Err.flush();
    return 0;
  }

  if (FLAGS_benchmark_list_tests) {
    for (auto const& benchmark : benchmarks) {
      Out << benchmark.name().str() << "\n";
    }
  } else {
    internal::RunBenchmarks(benchmarks, display_reporter, file_reporter);
  }
}
```

find benchmarks

returns a

vector of

BenchmarkInstance

After all BenchmarkInstances are found, creating a runner for each instance is the next stage.

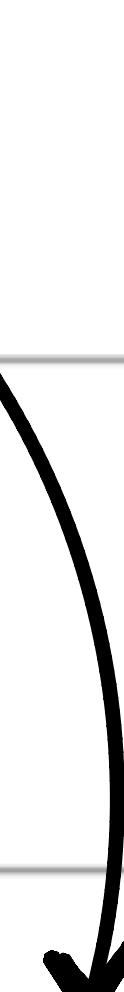
It is important to note that

a single Benchmark may

create more than one

BenchmarkInstance, check here:

```
// Loop through all benchmarks
for (const BenchmarkInstance& benchmark : benchmarks) {
    BenchmarkReporter::PerFamilyRunReports* reports_for_family = nullptr;
    if (benchmark.complexity() != oNone) {
        reports_for_family = &per_family_reports[benchmark.family_index()];
    }
    benchmarks_with_threads += static_cast<int>(benchmark.threads() > 1);
    runners.emplace_back(benchmark, &perfcounters, reports_for_family);
    int num_repeats_of_this_instance = runners.back().GetNumRepeats();
    num_repetitions_total +=
        static_cast<size_t>(num_repeats_of_this_instance);
    if (reports_for_family != nullptr) {
        reports_for_family->num_runs_total += num_repeats_of_this_instance;
    }
}
assert(runners.size() == benchmarks.size() && "Unexpected runner count.");
```



The next stage is creating a BenchmarkRunner for each instance.

Each runner will have to repeat R times. So, the next step is creating a data structure to ensure that a runner is repeating R times.

```
std::vector<size_t> repetition_indices;
repetition_indices.reserve(num_repetitions_total);
for (size_t runner_index = 0, num_runners = runners.size();
     runner_index != num_runners; ++runner_index) {
    const internal::BenchmarkRunner& runner = runners[runner_index];
    std::fill_n(first: std::back_inserter(&repetition_indices),
               n: runner.GetNumRepeats(), runner_index);
}
assert(repetition_indices.size() == num_repetitions_total &&
       "Unexpected number of repetition indexes.");
```

Create

repetition-indices

See here:

An instance of Runner has repetitions defined by GetNum.

- Repeate, which comes from BenchmarkInstance, which gets from Benchmark.

By default, repetition = 1.

There are two ways to define repetitions:

1) while defining a benchmark with BENCHMARK macro.

2) As command line argument

```
BenchmarkRunner::BenchmarkRunner(  
    BenchmarkReporter::PerFamilyRunReports* reports_for_family_)  
    : b_(b_),  
    reports_for_family(reports_for_family_),  
    parsed_benchtime_flag(ParseBenchMinTime(FLAGS_benchmark_min_time)),  
    min_time(FLAGS_benchmark_dry_run  
        ? 0  
        : ComputeMinTime(b_, parsed_benchtime_flag)),  
    min_warmup_time(  
        FLAGS_benchmark_dry_run  
        ? 0  
        : ((!IsZero(n: b_.min_time()) && b_.min_warmup_time() > 0.0)  
            ? b_.min_warmup_time()  
            : FLAGS_benchmark_min_warmup_time)),  
    warmup_done(FLAGS_benchmark_dry_run ? true : !(min_warmup_time > 0.0)),  
    repeats(FLAGS_benchmark_dry_run  
        ? 1  
        : (b_.repetitions() != 0 ? b_.repetitions()  
            : FLAGS_benchmark_repetitions)),
```

*repeate of Benchmarkrunner
comes from BenchmarkInstance
or from FLAGS_benchmark_repetitions*

FLAGS_benchmark_repetitions

```
// The number of runs of each benchmark. If greater than 1, the mean and  
// standard deviation of the runs will be reported.
```

```
BM_DEFINE_int32(benchmark_repetitions, 1);
```

↓ ↗ default value
name

Google benchmark creates some macros which are used to create command line flags.

Above image uses the macro: BM_DEFINE_int32 to create a variable of type int32_t with name as FLAGS_benchmark.

repetitions. This variable is assigned a function which reads

a command line flag:

benchmark_repetitions.

```
#define BM_DEFINE_int32(name, default_val) \  
BENCHMARK_EXPORT int32_t FLAG(name) = \  
benchmark::Int32FromEnv(#name, default_val)
```

```
#define FLAG(name) FLAGS_##name
```

Variable name →

FLAGS-benchmark-

repetitions.

```
BENCHMARK_EXPORT  
int32_t Int32FromEnv(const char* flag, int32_t default_val) {  
    const std::string env_var = FlagToEnvVar(flag);  
    const char* const value_str = getenv(name: env_var.c_str());  
    int32_t value = default_val;  
    if (value_str == nullptr ||  
        !ParseInt32(src_text: std::string(s: "Environment variable ") + env_var, value_str,  
                    &value)) {  
        return default_val;  
    }  
    return value;  
}
```

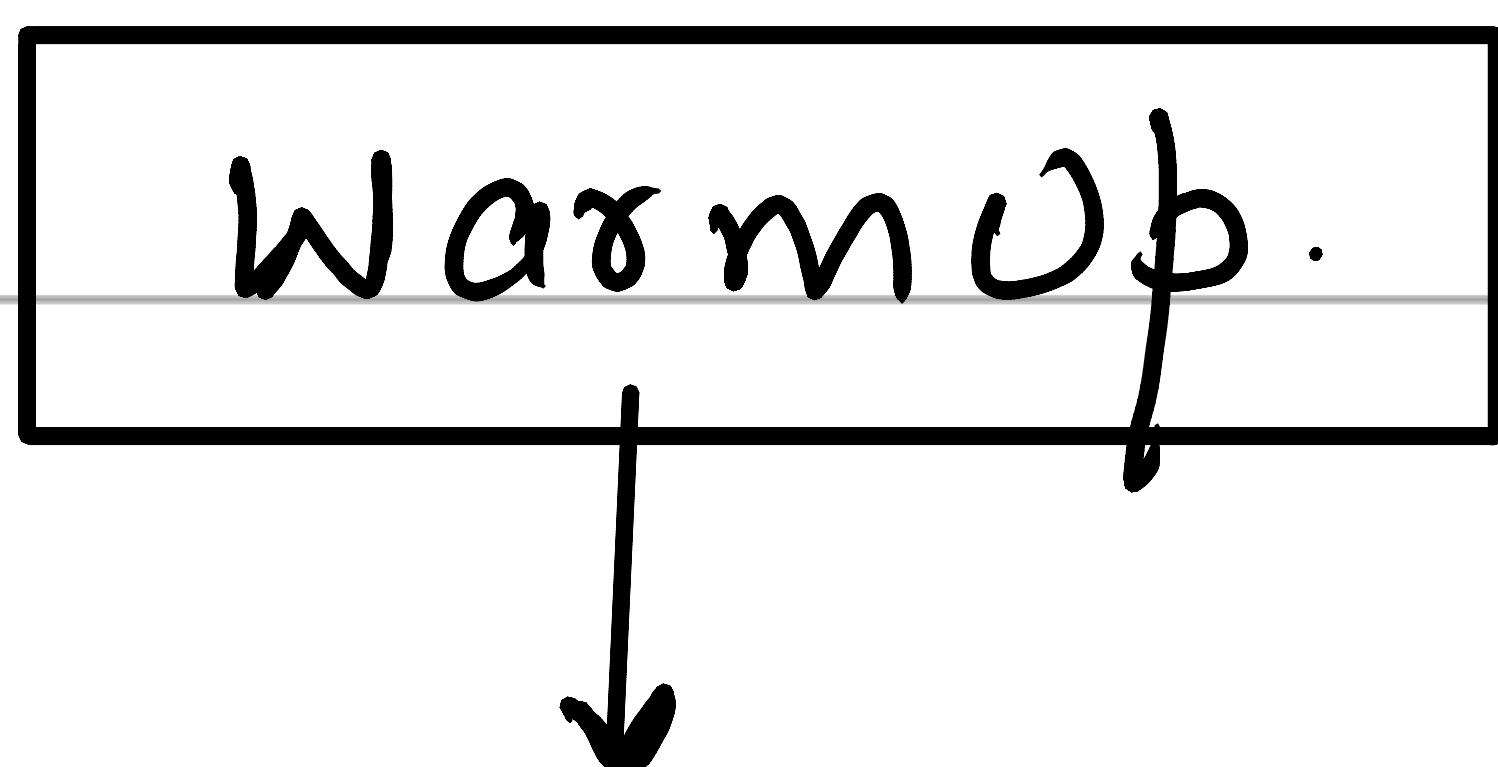
Function

to read the

command line argument, &
parse to int32.

Coming back to benchmark execution, after repetition-indices are created, each runner needs to do its repetitions.

Each repetitions involves :



Infinite loop of iterations



each iteration:

Setup, N iterations, Teardown

```
for (size_t repetition_index : repetition_indices) {  
    internal::BenchmarkRunner& runner = runners[repetition_index];  
    runner.DoOneRepetition();  
    if (runner.HasRepeatsRemaining()) {  
        continue;  
    }
```

One
Repetition

each runner

undergoes ~~its~~
repetitions.

```
void BenchmarkRunner::DoOneRepetition() {  
    // simply use that precomputed iteration count.  
    for (;;) {  
        b.Setup();  
        i = DoNIterations();  
        b.TearDown();  
  
        // Do we consider the results to be significant?  
        // If we are doing repetitions, and the first repetition was already done,  
        // it has calculated the correct iteration time, so we have run that very  
        // iteration count just now. No need to calculate anything. Just report.  
        // Else, the normal rules apply.  
        const bool results_are_significant = !is_the_first_repetition ||  
            has_explicit_iteration_count ||  
            ShouldReportIterationResults(i);  
  
        // Good, let's report them!  
        if (results_are_significant) {  
            break;  
        }  
  
        // Nope, bad iteration. Let's re-estimate the hopefully-sufficient  
        // iteration count, and run the benchmark again...  
  
        iters = PredictNumItersNeeded(i);  
        assert(iters > i.iters &&  
            "if we did more iterations than we want to do the next time, "  
            "then we should have accepted the current iteration run.");  
    }  
}
```

Warm Up

is done

before the

for loop.

Every time the loop runs, Setup
N iterations, and teardown are
called.

If enough iterations are done, the loop breaks, else the system predicts next iterations.

```
bool BenchmarkRunner::ShouldReportIterationResults(  
    const IterationResults& i) const {  
    // Determine if this run should be reported;  
    // Either it has run for a sufficient amount of time  
    // or because an error was reported.  
    return (i.results.skipped_ != 0u) || FLAGS_benchmark_dry_run ||  
        i.iters >= kMaxIterations || // Too many iterations already.  
        i.seconds >=  
            GetMinTimeToApply() || // The elapsed time is large enough.  
            // CPU time is specified but the elapsed real time greatly exceeds  
            // the minimum time.  
            // Note that user provided timers are except from this test.  
            ((i.results.real_time_used >= 5 * GetMinTimeToApply()) &&  
            !b.use_manual_time());  
}
```

enough iterations done?

(for now)

1) iteration

time \geq min bench time.

2) iterations \geq max supported iterations